

# Assignment 4

(Scope: Chapter 5 Process Scheduling)

Due Date: Sunday, October 28, 2018, 11:59pm

Submit electronically on iLMS

What to submit: One zip file named <studentID>-hw4.zip (replace <studentID> with your own student ID). It should contain seven files:

- one PDF file named **hw4.pdf** for Section 1 and Section 2. Write your answers in English. Check your spelling and grammar. Include your name and student ID!
- Section 2: Python source files. Include your name and student ID in the program comments on top.
  - Section 2.2: **task.py**
  - Section 2.3: **npsched.py** (for non-preemptive scheduler) and **typescript3**
  - Section 2.4 **psched.py** (for preemptive scheduler) and **typescript4**
  - Section 2.5 **typescript5**

## 1. [40 points] Problem Set

1. [20 points] 5.10 Which of the following scheduling algorithms could result in starvation? Explain.
  - a. First-come, first-served : No, because everyone has chance to do task in order.
  - b. Shortest job first: Yes, if there are always some processes are faster than the longest, the longest won't have a chance.
  - c. Round robin: No, because everyone has chance to do task within the time.
  - d. Priority :Yes, if there some priority of processes are higher than the longest, the longest won't have a chance.

Ans: b, d


2. [10 points] 5.12 Consider a system running ten I/O-bound tasks and one CPU-bound task. Assume that the I/O-bound tasks issue an I/O operation once for every millisecond of CPU computing and that each I/O operation takes 10 milliseconds to complete. Also assume that the context-switching overhead is 0.1 millisecond and that all processes are long-running tasks. Describe the CPU utilization for a round-robin scheduler when:
  - a. The time quantum is 1 millisecond
  - b. The time quantum is 10 milliseconds

(a) for every time quantum of a round-robin task is 1 millisecond and the scheduler should take 0.1 millisecond for context-switching for every process. Then CPU utilization is  $\frac{1}{1.1}$  (total time with context-switching) = 91%.

(b) for second problem, The time quantum is 10 milliseconds for ten I/O-bound tasks, they also have to take 1 millisecond to do context-switching, then for I/O-bound tasks is  $\frac{1.1 \times 10}{11} = 1$ ; for CPU-bound task executes for 10 milliseconds and 1 milliseconds for doing


context-switching which takes  $10+0.1 = 10.1$ . Hence, total time for CPU utilization is (I/O-bound tasks + CPU-bound task)/ total time. It is  $20/(11+10.1) = 20/21.1 = 94\%$ .

3. [10 points] Consider two processes,  $P_1$  and  $P_2$ , where  $p_1 = 50$ ,  $t_1 = 25$ ,  $p_2 = 75$ , and  $t_2 = 30$ .
- Can these two processes be scheduled using rate-monotonic scheduling? Illustrate your answer using a Gantt chart such as the ones in Figure 5.16-5.19.
  - Illustrate the scheduling of these two processes using earliest-deadline-first (EDF) scheduling.



Because  $P_1$  is faster than  $P_2$ , so priority is  $P_1 > P_2$ .  $P_1$  runs CPU burst at 25 time unit,  $P_2$  run till 50 time unit, But  $P_2$  remains 5 time units in first period, there is a miss. So the answer is no.

b.



Because  $P_1$  is faster than  $P_2$ , so priority is  $P_1 > P_2$ .  $P_1$  completes its CPU burst at 25 time unit, then  $P_2$  run till time unit 50, but within  $P_2$  deadline is at time 75, and  $P_1$  deadline is at time 100. so  $P_2$  has the earliest deadline so it completes its burst time after first period of  $P_1$ . At time 75,  $P_1$  runs first for the earliest deadline. Then  $P_2$  starts to run till time 100. At time 150,  $P_1$  and  $P_2$  have the same deadline, so  $P_2$  continues bursting time after  $P_1$  complete its burst time.

## 2. [60 points] Programming Exercise

In this programming exercise, you are to build a CPU scheduler that can compute the schedule for a variety of policies and calculate the various cost functions.

## 2.1 FIFO and Priority Queue

A fundamental data structure in any CPU scheduler is a queue. Here, it can refer to a FIFO (first-in first-out) queue, but it may also refer to a priority queue, a LIFO (last-in first-out, also known as a stack), etc. Unlike random-access memory, where the reader or writer provides the memory address explicitly, a queue keeps track of its own addresses and provides only `.get()` and `.put()` methods for reading and writing one element at a time. The following class is provided as an example:

----- file "[fifo.py](#)" -----

```
class FIFO:
    def __init__(self, initList=[]):
        self.A = list(initList)
    def get(self):          # remove element and return its value
        return self.A.pop(0) # throws underflow exception if empty
    def put(self, val):     # add element
        self.A.append(val)
    def head(self):        # A[0] if not empty, None instead of underflow exception
        return len(self.A) and self.A[0] or None
    def __iter__(self):    # iterator over its elements
        for i in self.A:  # convertible to tuple, list, for-in loop, etc
            yield i
    def __len__(self):     # allows caller to call len(f) where f is FIFO
        return len(self.A)
    def __repr__(self):   # shows a representation; we just show it as list
        return repr(self.A)
```

This will handle any data type. An example is (assume you save it in `fifo.py`)

```
>>> from fifo import *
>>> f = FIFO(range(3))
>>> f
[0, 1, 2]
>>> f.put(6)
>>> f.get()
0
>>> f.head()
1
>>> len(f)
3
```

In addition, we also provide an implementation of a priority queue based on min-heap. You don't get the source code but you can import it from the `minheap.pyc` file provided:

[minheap.pyc](#) (for python 3.5.2 -- use this); or [minheap.pyc](#) (for python 2.7.13)

[minheap.pyc](#) (for python 3.5.3); or [minheap.pyc](#) (for python 3.6.2)

It has the following API:

----- file "[minheap.py](#)" -----

```
class MinHeap:
    def __init__(self):
    def __len__(self):
    def __iter__(self):
    def __repr__(self):
    def get(self):
    def put(self, value):
```

```

def head(self):
def buildheap(self): # reinitialize content to be heap again

```

One difference is that your minheap data structure typecasts its elements to tuples before comparison, and Python will compare tuples in lexicographical order, and we will exploit this characteristic later when prioritizing tasks to run.

```

>>> from minheap import MinHeap
>>> h = MinHeap()
>>> for i in [(2,3), (3,4), (2,4), (4,5), (5, 6)]: h.put(i)
...
>>> h
[(2, 3), (3, 4), (2, 4), (4, 5), (5, 6)]
>>> h.get()
(2, 3)
>>> h
[(2, 4), (3, 4), (5, 6), (4, 5)]
>>> h.get()
(2, 4)
>>> h
[(3, 4), (4, 5), (5, 6)]
>>> h.put((6,7))
>>> h.get()
(3, 4)
>>> h
[(5, 6), (4, 5), (6, 7)]

```

## 2.2 Task class [10 points]

You need to declare a `Task` class for representing the properties of a task to be scheduled, including properties given by the user and additional data for bookkeeping purpose. Here, we use the term `Task` to mean the workload to be performed, with or without having a process or a thread attached to it. A thread or process may be recycled to run different tasks over time. But sometimes tasks and processes are used interchangeably when the task is attached to a process. The given data are passed as arguments to the constructors. You may use the following template to define your task. Look for the italicized comments to add your own code.

----- file "[task-template.py](#)" : save and rename it as "task.py" -----

```

class Task:
    def __init__(self, name, release, cpuBurst):
        # the task has a string name, release time and cpuBurst.
        # the constructor may also need to initialize other fields,
        # for statistics purpose. Examples include
        # waiting time, remaining time, last dispatched time, and
        # completion time

    def __repr__(self):
        # returns a string that looks like constructor syntax
        return self.__class__.__name__ + '(%s, %d, %d)' % (repr(self.name), \
            self.release, self.cpuBurst)

    def __str__(self):
        return self.name

    _KNOWN_SCHEMES = ["FCFS", "SJF", "RR"]
    def setPriorityScheme(self, scheme="SJF"):
        """

```

```

        the scheme can be "FCFS", "SJF", "RR", etc
    """
    self.scheme = scheme
    if not scheme in _KNOWN_SCHEMES:
        raise ValueError("unknown scheme %s: must be FCFS, SJF, RR" %
scheme)
    def decrRemaining(self):
        # call this method to decrement the remaining time of this task
    def remainingTime(self):
        # returns the remaining time of this task
    def done(self):
        # returns a boolean for if this task has remaining work to do
    def setCompletionTime(self, time):
        # record the clock value when the task is completed
    def turnaroundTime(self):
        # returns the turnaround time of this task, as defined on
        # slide 9 of chapter-5 lecture
    def incrWaitTime(self):
        # increment the amount of waiting time for this task
    def releaseTime(self):
        # returns the release time of this task
    def __iter__(self):
        # this enables converting the task into a tuple() type so that
        # the priority queue can just cast it to tuple before comparison.
        # it depends on the policy
        if (self.scheme == 'FCFS'):
            t = # your tuple that defines the priority
        elif (self.scheme == 'SJF'): # shortest job first
            t = # your tuple that defines the priority
        elif (self.scheme == 'RR'): # round robin
            t = # your tuple that defines the priority
        else:
            raise ValueError("Unknown scheme %s" % self.scheme)
        for i in t:
            yield i

```

## 2.3 Nonpreemptive Scheduler (20 points)

The NPScheduler class is instantiated with a policy and up to N time steps. Then the caller may add tasks to be scheduled, either as the scheduler runs or all at the beginning. The scheduler runs one time step at a time to fill in the Gantt chart with scheduled tasks. It also provides methods for the statistics. Use the following template ([npsched-template.py](#), rename it as `npsched.py`) to make your scheduler

```

from fifo import FIFO
from minheap import MinHeap
from task import Task
class NPScheduler: # nonpreemptive scheduler
    def __init__(self, N, policy='SJF'):
        self.N = N # number of timesteps to schedule
        self.running = None
        self.clock = 0 # the current timestep being scheduled
        self.policy = policy

```

```

    # instantiate the readyQueue, which may be a FIFO or MinHeap
    # you may need additional queues for
    # - tasks that have been added but not released yet
    # - tasks that have been completed
    # - the Gantt chart

def addTask(self, task):
    # if the release time of the new task is not in the future, then
    # put it in ready queue; otherwise, put into not-ready queue.
    # you may need to copy the scheduler policy into the task

def dispatch(self, task):
    # dispatch here means assign the chosen task as the one to run
    # in the current time step.
    # the task should be removed from ready-queue by caller;
    # The task may be empty (None).
    # This method will make an entry into the Gantt chart and perform
    # bookkeeping, including
    # - recording the last dispatched time of this task,
    # - increment the wait times of those tasks not scheduled
    #   but in the ready queue

def releaseTasks(self):
    """
    this is called at the beginning of scheduling each time step to see
    if new tasks became ready to be released to ready queue, when their
    release time is no later than the current clock.
    """
    while True:
        r = self.notReadyQueue.head()
        # assuming the not-Ready Queue outputs by release time
        if r is None or r.releaseTime() > self.clock:
            break
        r = self.notReadyQueue.get()
        r.setPriorityScheme(self.policy)
        self.readyQueue.put(r)

def checkTaskCompletion(self):
    # if there is a current running task, check if it has just finished.
    # (i.e., decrement remaining time and see if it has more work to do.
    # If so, perform bookkeeping for completing the task,
    # - move task to done-queue, set its completion time and lastrun time
    # set the scheduler running task to None, and return True
    # (so that a new task may be picked.)
    # but if not completed, return False.
    # If there is no current running task, also return True.
    if self.running is None:
        return True
    # your code here

def schedule(self):
    # scheduler that handles nonpreemptive scheduling.
    # the policy such as RR, SJF, or FCFS is handled by the task as it
    # defines the attribute to compare (in its __iter__() method)
    # first, check if added but unreleased tasks may now be released

```

```

# (i.e., added to ready queue)
self.releaseTasks()
if self.checkTaskCompletion() == False:
    # There is a current running task and it is not done yet!
    # the same task will continue running to its completion.
    # simply redispach the current running task.
else:
    # task completed or no running task.
    # get the next task from priority queue and dispatch it.

def clockGen(self):
    # this method runs the scheduler one time step at a time.
    for self.clock in range(self.N):
        # now run scheduler here
        self.schedule()
        yield self.clock

def getSchedule(self):
    return '-'.join(map(str, self.ganttChart))

def testNPScheduler(tasks, policy):
    nClocks = 20
    scheduler = NPScheduler(nClocks, policy)

    for t in tasks:
        scheduler.addTask(t)

    for clock in scheduler.clockGen():
        pass

    print('nonpreemptive %s: %s' % (scheduler.policy,
    scheduler.getSchedule()))

if __name__ == '__main__':
    tasks = [Task(*i) for i in [('A', 0, 7), ('B', 2, 4), ('C', 4, 1), ('D',
5, 4)]]
    print('tasks = %s' % tasks)
    for policy in ['SJF', 'FCFS', 'RR']:
        tasks = [Task(*i) for i in [('A', 0, 7), ('B', 2, 4), ('C', 4, 1),
('D', 5, 4)]]
        testNPScheduler(tasks, policy)

```

----- Your output would look like this:

```

$ python3 npscheduler.py
tasks = [Task('A', 0, 7), Task('B', 2, 4), Task('C', 4, 1), Task('D', 5,
4)]
nonpreemptive SJF: A-A-A-A-A-A-A-C-B-B-B-B-D-D-D-D-None-None-None-None
nonpreemptive FCFS: A-A-A-A-A-A-A-B-B-B-B-B-C-D-D-D-D-None-None-None-None
nonpreemptive RR: A-A-A-A-A-A-A-B-B-B-B-B-C-D-D-D-D-None-None-None-None

```

## 2.4 Preemptive Scheduler (20 points)

For this part, make a copy of your nonpreemptive scheduler and make it a preemptive one.

The overall structure is the same as the Nonpreemptive scheduler.

----- file "[psched-template.py](#)", rename and save as "psched.py"

```
class PScheduler(NPScheduler): # subclass from nonpreemptive scheduler
    # this means it can inherit
    # __init__(), addTask(), dispatch(), releaseTasks()
    # clockGen(), getSchedule()

    def preempt(self):
        # this is the new method to add to put the running task
        # back into ready queue, plus any bookkeeping if necessary.

    def schedule(self):
        self.releaseTasks() # same as before
        if self.checkTaskCompletion() == False:
            # still have operation to do.
            # see if running task or next ready task has higher priority
            # hint: compare by first typecasting the tasks to tuple() first
            # and compare them as tuples. The tuples are defined in
            # the __iter__() method of the Task class based on policy.
            # if next ready is not higher priority, redispach current task.
            # otherwise,
            # - swap out current running (by calling preempt method)
            # task completed or swapped out
            # pick next task from ready queue to dispatch, if one exists.

def testPScheduler(tasks, policy):
    # this is same as before, but instantiate the preemptive scheduler.
    nClocks = 20
    scheduler = PScheduler(nClocks, policy)
    # the rest is the same as before
    for t in tasks:
        scheduler.addTask(t)
    for clock in scheduler.clockGen():
        pass
    print('preemptive %s: %s' % (scheduler.policy, scheduler.getSchedule()))

if __name__ == '__main__':
    tasks = [Task(*i) for i in [('A', 0, 7), ('B', 2, 4), ('C', 4, 1), ('D',
5, 4)]]
    print('tasks = %s' % tasks)
    for policy in ['SJF', 'FCFS', 'RR']:
        tasks = [Task(*i) for i in [('A', 0, 7), ('B', 2, 4), ('C', 4, 1),
('D', 5, 4)]]
        testPScheduler(tasks, policy)
```

Your output would look like

```
tasks = [Task('A', 0, 7), Task('B', 2, 4), Task('C', 4, 1), Task('D', 5,
4)]
preemptive SJF: A-A-B-B-C-B-B-D-D-D-D-A-A-A-A-A-None-None-None-None
```



```
preemptive FCFS: A-A-A-A-A-A-A-B-B-B-B-C-D-D-D-D-None-None-None-None
preemptive RR: A-A-B-A-B-C-A-D-B-A-D-B-A-D-A-D-None-None-None-None
```

## 2.5 Add Statistics (10 points)

Implement the following methods to the nonpreemptive scheduler code (and the preemptive one will automatically get the same code due to inheritance).

```
def getThroughput(self):
    # throughput is the number of processes completed per unit time.
    # returns a tuple for (number of done processes, number of clocks)

def getWaitTime(self):
    # returns a tuple for (total wait time of processes, #processes)

def getTurnaroundTime(self):
    # returns a tuple for (total turnaround times, #processes)
```

Combine the nonpreemptive and preemptive schedulers into the same test bench and print out the statistics. Download the [schedstat.py](#) to run, and the output looks like

```
$ python3 schedstat.py
tasks = [Task('A', 0, 7), Task('B', 2, 4), Task('C', 4, 1), Task('D', 5, 4)]
nonpreemptive SJF: A-A-A-A-A-A-A-C-B-B-B-B-D-D-D-D-None-None-None-None
    thruput = (4, 16) = 0.25, waittimes = (16, 4) = 4.00, turnaroundtime = (32, 4) = 8.00
preemptive SJF: A-A-B-B-C-B-B-D-D-D-D-A-A-A-A-A-None-None-None-None
    thruput = (4, 16) = 0.25, waittimes = (12, 4) = 3.00, turnaroundtime = (28, 4) = 7.00
nonpreemptive FCFS: A-A-A-A-A-A-A-B-B-B-B-C-D-D-D-D-None-None-None-None
    thruput = (4, 16) = 0.25, waittimes = (19, 4) = 4.75, turnaroundtime = (35, 4) = 8.75
preemptive FCFS: A-A-A-A-A-A-A-B-B-B-B-C-D-D-D-D-None-None-None-None
    thruput = (4, 16) = 0.25, waittimes = (19, 4) = 4.75, turnaroundtime = (35, 4) = 8.75
nonpreemptive RR: A-A-A-A-A-A-A-B-B-B-B-C-D-D-D-D-None-None-None-None
    thruput = (4, 16) = 0.25, waittimes = (19, 4) = 4.75, turnaroundtime = (35, 4) = 8.75
preemptive RR: A-A-B-A-B-C-A-D-B-A-D-B-A-D-A-D-None-None-None-None
    thruput = (4, 16) = 0.25, waittimes = (22, 4) = 5.50, turnaroundtime = (38, 4) = 9.50
```